

CSCE 878 (Fall 2016) Homework 1

Zeynep Mutlu Hakguder

Department of Electrical and Computer Engineering

University of Nebraska-Lincoln, USA

zeynep.hakguder@huskers.unl.edu

Haluk Dogan

Department of Electrical and Computer Engineering

University of Nebraska-Lincoln, USA

haluk.dogan@huskers.unl.edu

September 26, 2016

Contents

1	Question 1	3
2	Question 2	5
3	Question 3	6
4	Question 4 and 5	7
4.1	Introduction	7
4.2	Data	8
4.3	Implementation	8
4.4	Results	10
4.5	Discussion	11

1 Question 1

The training data consists of 6 examples each with a single integer-valued attribute and a binary label. The target function \mathcal{C} is represented as a single interval using two points a and b .

1	-	6		11	
2		7	+	12	-
3		8	+	13	
4		9		14	
5	+	10		15	-

Table 1: Training Set

- (a) A hypothesis consistent with the training set would be obtained if $a = 5$ and $b = 8$ so that an instance x is labeled positive if and only if $5 \leq x \leq 8$.
- (b) The version space is the subset of hypothesis space \mathcal{H} (which is chosen as the family of single intervals and is guaranteed to include target concept) that is consistent with the training set. For example, in the interval given in (a), all positive training examples are contained, if it was any smaller, it would exclude a positive training example and be inconsistent with the data. We could expand the interval to $2 \leq x \leq 11$, this is the largest interval that is consistent with the training set, if we were to expand it any further we would end up including a negative example in the interval that represents the target function. The version space is the set of all such hypotheses (represented by intervals) that are consistent with the training set. The interval's lower bound can be any of 2, 3, 4 or 5, its upper bound can be any of 8, 9, 10 or 11, the size of the version space is $4 \times 4 = 16$.
- (c) To decrease the size of the version space we can specify the query $q_1 = 10$. If the label for 10 is:
- positive, the smallest interval consistent with our training data and q_1 would be $5 \leq x \leq 10$. We would eliminate 8 hypotheses that assign negative labels

to 9 and 10 from the version space.

- negative, the smallest interval consistent with our training data and q_1 would be $5 \leq x \leq 9$. We would eliminate 8 hypotheses that assign positive labels to 10 and 11 from the version space.

If we set $q_2 = 6$, this query wouldn't change the size of the version space regardless of the answer.

- (d) If we have a training set with three examples $\{(x_1, -), (x_2, +), (x_3, -)\}$, where $0 < x_1 < x_2 < x_3$, in order to reduce the size of the version space, we can form queries in a manner similar to binary search. First, to find the upper bound of the interval we ask the label of the midpoint between x_2 and x_3 , depending on the label we receive we form our next query; i.e if it is negatively labeled we ask the label for the midpoint between the last query and x_2 , if it is positively labeled we ask the label for the midpoint between the last query and x_3 . Similarly, to obtain the lower boundary we ask the label for the midpoint between x_1 and x_2 , if it is negatively labeled we ask the label for the midpoint between the last query and x_2 , if it's positively labeled we ask the label for the midpoint between the last query and x_1 . Below is an example showing one of the worst case scenarios with training set $\{(1, -), (9, +), (17, -)\}$.




Figure 1: Query example to decrease version space size.

2 Question 2




Figure 2: $A \wedge [\sim B]$

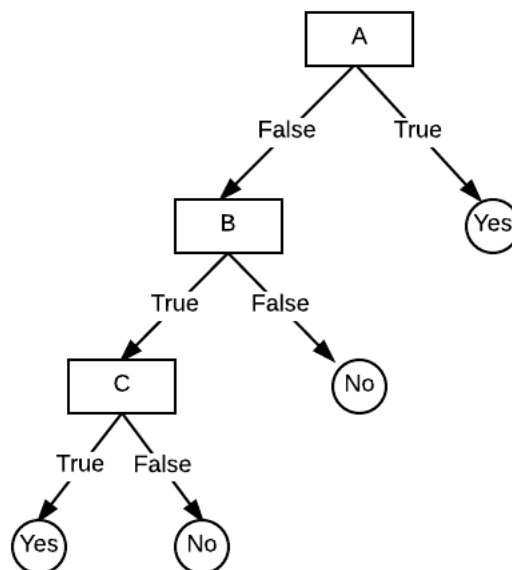


Figure 3: $A \vee [B \wedge C]$




Figure 4: $A \oplus B$

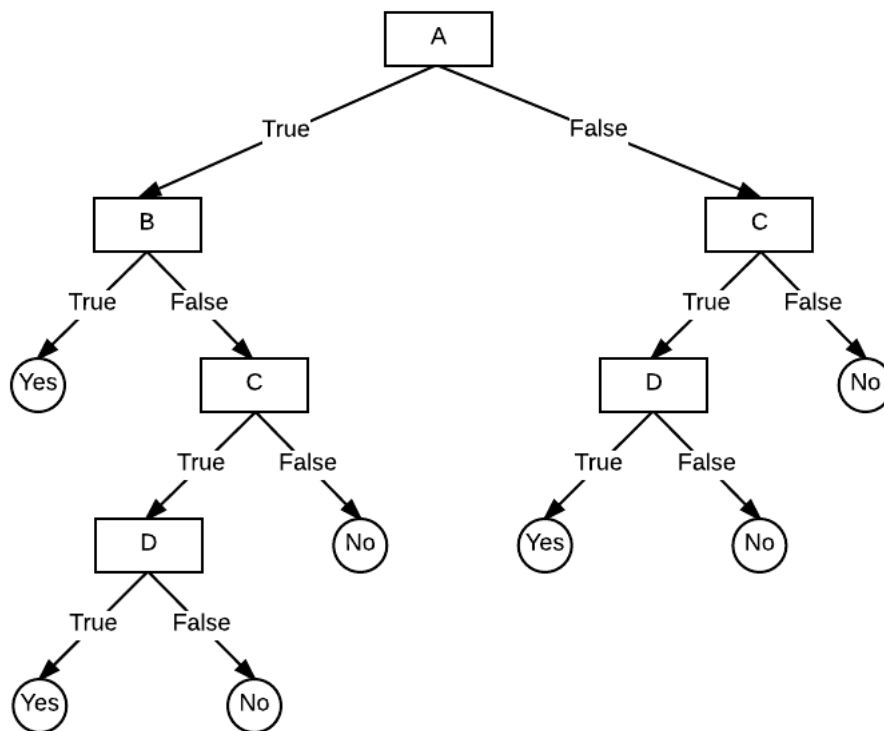


Figure 5: $[A \wedge B] \vee [C \wedge D]$

3 Question 3

- (a) The entropy of the dataset is calculated as follows:

$$\begin{aligned}
I &= -p^+ \log_2 p^+ - p^- \log_2 p^- \\
&= -\frac{3}{6} \log_2 \frac{3}{6} - \frac{3}{6} \log_2 \frac{3}{6} \\
&= 1
\end{aligned}$$

(b) The ID3 algorithm would choose the attribute with smallest impurity which is calculated as below:

$$\begin{aligned}
I'(a_i) &= \frac{|T|}{|T+F|} (-p_T^+ \log_2 p_T^+ - p_T^- \log_2 p_T^-) \\
&\quad + \frac{|F|}{|T+F|} (-p_F^+ \log_2 p_F^+ - p_F^- \log_2 p_F^-)
\end{aligned}$$

where $i = 1, 2$.

$$\begin{aligned}
I'(a_1) &= \frac{3}{6} \left(-\frac{2}{3} \log_2 \frac{2}{3} - \frac{1}{3} \log_2 \frac{1}{3} \right) + \frac{3}{6} \left(-\frac{1}{3} \log_2 \frac{1}{3} - \frac{2}{3} \log_2 \frac{2}{3} \right) = 0.9183 \\
I'(a_2) &= \frac{4}{6} \left(-\frac{2}{4} \log_2 \frac{2}{4} - \frac{2}{4} \log_2 \frac{2}{4} \right) + \frac{2}{6} \left(-\frac{1}{2} \log_2 \frac{1}{2} - \frac{1}{2} \log_2 \frac{1}{2} \right) = 1.0
\end{aligned}$$

Since $I'(a_1) < I'(a_2)$, ID3 algorithm would choose a_1 next.

4 Question 4 and 5

4.1 Introduction

We've implemented the ID3 algorithm with impurity as splitting criterion along with rule post-pruning. We used 4 data sets from the UC Irvine Machine Learning Repository. In this report, we present the results of the ID3 algorithm on these data sets.

4.2 Data

4 data sets were obtained from UC Irvine Machine Learning Repository. In addition to the “Congressional Voting Records (house-voting-84)” and “Monks I” data sets, we chose “Nursery” and “Tic-Tac-Toe Endgame” datasets. All of these data sets have categorical attribute values with number of attributes ranging from 7 to 16. All except the “Nursery” data set have binary class labels, “Nursery” data set has 5 possible class values. None of the data sets have missing attribute values except “Congressional Voting Records”. Attributes in “Congressional Voting” data set take “?” as values, however we considered this a legitimate attribute value carrying information given the context. In addition, we used the weather data set used in lectures to verify our implementation. All data sets can be found in the accompanying archive. Summary of the data sets can be found in the Table 2.

Data Set Name	house-votes-84	monks-1	tic-tac-toe	nursery
Data Set Characteristics	Multivariate	Multivariate	Multivariate	Multivariate
Number of Instances	434	556	958	12960
Number of Attributes	16	7	9	8
Attribute Characteristics	Categorical	Categorical	Categorical	Categorical
Missing Values?	No	No	No	No

Table 2: Summary of the selected data sets.

4.3 Implementation

We implemented ID3 algorithm with post-pruning with Scala (source code available in archive). For each data set, we, first, reserved more than 30 examples sampled randomly without replacement for testing, then from the remaining training set we extracted at least 30 examples randomly without replacement for validation. We take these values as inputs to the “.properties” file. After building the tree, we extracted the rule set by traversing from the root to each leaf. After obtaining each leaf, we used the validation set to prune these rules. We dropped a precedent in a rule, if when dropped the modified rule had better accuracy on the validation set. We sorted the rules according to their

validation accuracy in a decreasing manner.

While predicting the classes of the examples in the test set, we traversed the tree with each example and whenever an example had an attribute value unmatched in the tree we used majority voting (see 4.5). Prediction with pruned rule set was done by checking each example against the rules in the rule set beginning with the rule with highest accuracy. When an example had matching attribute values with a rule, we used that rule to predict the class, and continued with the next example.

After running the program the following files can be found:

- “.validation” contains the examples that were used for pruning the tree.
- “.test” contains the examples that the tree and the pruned rule set were tested on.
- “.training” contains the examples left in the training set after test and validation sets are formed.
- “.tree” contains the tree built by ID3 algorithm, each tab signifies added depth to the tree with a chosen attribute. The attribute used to split the data is indicated by ”Decide-on”, entropy of the dataset at that node and the impurity of the chosen attribute as well as the number of class members at that node are reported.
- “.prunedRuleSetSorted” gives the accuracy of each pruned rule on the validation set in sorted order by accuracy
- “.prunedRuleSetUnsorted” gives the accuracy of each pruned rule on the validation set
- “.prunedRuleSetTest” gives the accuracy of each pruned rule on the test set
- “.unprunedRuleSet” gives the origin rule set extracted from the built tree
- “.unprunedRuleSetTest” gives the test accuracies of unpruned rules

4.4 Results

We observed comparable heights obtained for different datasets. Trees had heights ranging between 5 and 8 on average over 100 iterations even though the attribute values change between 7 and 16. However, the number of nodes ranged between 21 and 320 on average over 100 iterations (Table 2, Table 3).

Data Set Name	house-votes-84	monks-1	tic-tac-toe	nursery
Avg. number of nodes	21.01	38.95	112.51	319.79
Avg. height	7.23	5.21	7.33	8.0
Avg. accuracy of tree on test data	0.936	0.966	0.833	0.983
Avg. accuracy of pruned rules on test data	0.943	0.984	0.861	0.996

Table 3: Results of 100 iterations (Validation size = 30, Test size = 30).

We always observed higher accuracies for the pruned rule set than the tree. This means that overfitting took place while building the tree, i.e we can find an alternative hypothesis with higher training but lower test error. Pruning the rules improved accuracy (Figures 6, Figure 7, and Figure 8).




Figure 6: Accuracy comparison of tree and post pruning on house-votes-84 test data with changing training data size.




Figure 7: Accuracy comparison of tree and post pruning on monks-1 test data with changing training data size.




Figure 8: Accuracy comparison of tree and post pruning on tic-tac-toe test data with changing training data size.

4.5 Discussion

We observed comparable tree heights and different node numbers for different datasets. All trees except for “Congressional Voting” have heights that are about the same as their attribute numbers. This would mean for the “Congressional Voting” data set, that 7 attributes are enough to classify an example despite the larger available attribute number. The number of nodes in “Tic-Tac-Toe” and “Nursery” data sets were much larger than in “Monks-I” and “Congressional Votes” data sets, this is commensurate with the number

of available values for attributes in each data set (Table 2, Table 3).

During implementation, we were faced with a number of issues. First, during prediction of novel examples using the tree, some of the attributes had values in test data that were not present in our tree. For example, while our tree could split the data into “hot” and “cool” values of temperature attribute, it couldn’t decide which subset an example with “mild” temperature value belonged to (depicted as in Figure 9) . In that case, we opted for sending the example as deep in the tree as possible, then classifying the example as a member of the group with highest frequency in that node. If there were more than one classes with the highest frequency we picked the class randomly.

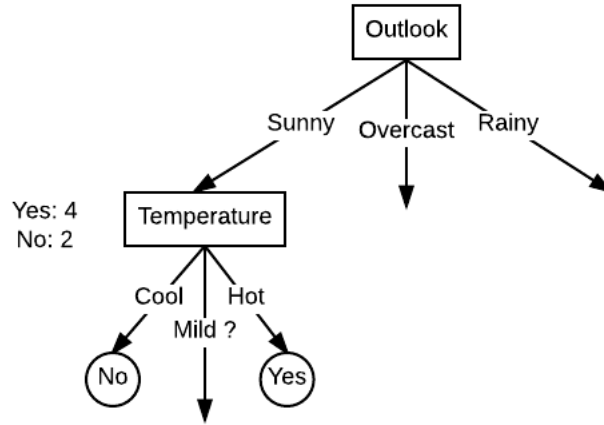


Figure 9: Depiction of classifying the example as a member of the group with the highest frequency.

During pruning, there was another issue we had to make a implementation decision upon. After building the tree and obtaining the rules, some of the rules applied to none of the examples in the validation set. Similarly, we observed the same issue even with rules with reduced precedents. In those cases, we dropped a precedent only if the modified rule (rule with the precedent dropped) applied to an example and it’s accuracy was at least 0.9. In other cases, where the rule (with the precedent we were considering dropping) applied to some examples, we only dropped the precedent if the modified rule had higher accuracy than the unmodified rule.

To gauge overfitting, we decreased the size of the training set. We observed a steady

decrease in the average accuracy of both the tree and the pruned rule set in “Monks-1” data, although the decrease was more obvious in the tree predictions as depicted in Figure 7. This could mean that overfitting occurs when the data is not sufficiently large because the algorithm is not presented with enough examples to detect true rules. However, this was not the case for other datasets with comparable sizes.